

Observability-Driven Intelligence for Anomaly Detection in Cloud-Native Distributed Microservices

Ngoc Duc Anh Nguyen ✉

School of Built Environment, Engineering and Computing, Leeds Beckett University, Leeds, UK

Satish Kumar ✉

School of Built Environment, Engineering and Computing, Leeds Beckett University, Leeds, UK

Nawar Jawad ✉

School of Built Environment, Engineering and Computing, Leeds Beckett University, Leeds, UK

Abstract

Automated fault diagnosis in cloud-native microservice architectures remains an open challenge: the three pillars of observability—metrics, logs, and distributed traces—each capture a distinct failure dimension, yet most production AIOps deployments exploit only one signal modality, leaving cross-modal causal chains undetected until user-visible degradation occurs. We present an observability-driven AIOps platform that closes this gap through continuous, fully offline multi-modal intelligence, motivated by AML compliance environments where data residency regulations categorically prohibit telemetry export to third-party infrastructure. Eight Kubernetes-native AIOps services ingest a three-pillar telemetry pipeline (Prometheus, Loki, Grafana Tempo) and apply a three-algorithm ML pipeline in sequence: Isolation Forest anomaly scoring over 60-second stream-windowed feature vectors, Pearson correlation root cause ranking over sliding observational windows, and ordinary least-squares regression on p95 latency trends for predictive SLO breach forecasting. Rather than delegating incident narration to a cloud-hosted model, the platform synthesises MELT context through a locally deployed Qwen2.5-3B large language model [27] served via Ollama [24] and continuously fine-tuned with LoRA adapters [16] on resolved incident outcomes—producing structured, actionable JSON diagnostic reports with no external API call leaving the cluster. The platform is evaluated through controlled Chaos Mesh [5] fault injection across four fault classes, addressing three research questions: whether fusing all three signal modalities outperforms any single-modal configuration (RQ1), whether correlation-based root cause ranking localises faults more accurately than threshold or trace-error baselines (RQ2), and what overhead full-fidelity distributed tracing imposes on regulated workloads (RQ3). Preliminary validation confirms fault detection within two 60-second analysis cycles, on-cluster LLM narration at 3–5 seconds per report, and zero telemetry egress over a 2-hour injection session—establishing the platform as both a privacy-preserving operational tool and a reproducible experimental substrate for quantitative AIOps evaluation.

2012 ACM Subject Classification Computer systems organization → Cloud computing; Computing methodologies → Anomaly detection; Information systems → Data management systems

Keywords and phrases anomaly detection, microservices, observability, machine learning, Kubernetes, distributed tracing, root cause analysis, AIOps, anti-money laundering

Digital Object Identifier 10.4230/LIPIcs...

Funding This research is supported by Leeds Beckett University.

1 Introduction

Distributed microservice architectures now coordinate hundreds of containerised services across Kubernetes clusters, each emitting metrics, logs, and distributed traces at sub-second resolution [4]. Anomalies in such systems manifest as subtle cross-service causal chains—a



© Ngoc Duc Anh Nguyen, Satish Kumar, and Nawar Jawad;
licensed under Creative Commons License CC-BY 4.0

Leibniz International Proceedings in Informatics

LIPICS Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

memory leak first inflects a metric, propagates to trace latency shifts, and eventually cascades into error-log floods—patterns invisible to any single-modal detector and intractable for human operators at the telemetry volumes modern clusters produce.

In regulated domains such as anti-money laundering (AML), this challenge is compounded by data residency obligations. Financial regulations in most jurisdictions explicitly prohibit exporting operational telemetry to third-party cloud services [18], categorically excluding commercially hosted AIOps platforms precisely in the environments where automated anomaly detection is most consequential. AML workloads further exhibit architectural heterogeneity—synchronous KYC REST APIs, asynchronous Kafka event streams, and stateful case-management workflows with legally mandated lifecycle invariants—that generates telemetry across fundamentally different operational patterns simultaneously.

Existing approaches are inadequate on two fronts. Rule-based monitors evaluate individual metrics against static thresholds, generating excessive false-positive alert volumes and failing on distributed cross-service anomalies [31]. Supervised learning approaches require labelled training data unavailable for novel failure modes. Commercial AIOps platforms—Dynatrace Davis AI [8] and Datadog Watchdog [7]—apply sophisticated correlation algorithms but mandate continuous telemetry export to vendor-managed infrastructure. No existing open-source platform unifies multi-modal anomaly detection, correlation-based RCA, and predictive SLO forecasting as a privacy-preserving, fully offline Kubernetes-native service.

We address this gap with two complementary design decisions: (i) all intelligence executes within the cluster—Isolation Forest [21] anomaly scoring, Pearson correlation root cause ranking, OLS SLO breach prediction, and Qwen2.5-3B [27] incident narration fine-tuned with LoRA [16] via Ollama [24], with no external API access; and (ii) the AML testbed uses a transactional outbox pattern that dual-routes domain events as both Kafka streams and OpenTelemetry span events carrying business attributes, providing naturally labelled ground-truth anomaly signals without secondary annotation.


The main contributions are: (i) a production-grade, open-source eleven-service Kubernetes platform for AML-domain observability with a fully offline intelligence layer satisfying financial data residency requirements; (ii) a domain event-based instrumentation strategy deriving ground-truth anomaly labels from AML compliance state transitions via OpenTelemetry span attributes; (iii) a three-algorithm offline ML pipeline (Isolation Forest, Pearson correlation RCA, OLS SLO prediction) integrated with a locally fine-tuned Qwen2.5-3B served via Ollama; and (iv) a controlled Chaos Mesh fault injection evaluation protocol addressing signal sufficiency (RQ1), ML-driven RCA localisation (RQ2), and distributed tracing overhead on regulated workloads (RQ3).

2 Background and Related Works

2.1 Cloud-Native Observability

Observability in distributed systems is the capacity to reconstruct application state from emitted telemetry without requiring instrumentation changes or direct access to production processes [1]. This distinguishes observability from monitoring: monitoring tests predefined conditions, whereas observability enables answering questions not anticipated at deployment time.

The operational realisation of observability rests on three interdependent signal types (Fig. 1). *Metrics* are scalar measurements sampled at regular intervals—CPU utilisation, request throughput, error rates—providing low-overhead, high-frequency health indicators; Prometheus is the de facto standard for Kubernetes metrics collection [2]. *Logs* are immutable



■ **Figure 1** Three-pillar cloud-native observability. Microservices auto-instrumented via OpenTelemetry emit metrics (Prometheus), logs (Loki), and distributed traces (Tempo). The intelligence layer fuses all three signal types for anomaly detection, RCA, and predictive forecasting.

timestamped records of discrete events; Loki extends the Prometheus model to log aggregation via LogQL [11]. *Distributed traces* follow individual requests across services, recording span durations, parent-child call relationships, and propagated context—providing the causal structure that neither metrics nor logs alone supply; Grafana Tempo provides a cost-effective trace backend for Kubernetes-scale ingestion [12].

The observability challenge is amplified in cloud-native architectures where cascading failures—a transient spike in one shared layer propagating through dozens of dependent services [31]—and ephemeral pod lifetimes make historical baselines a moving target. Sigelman et al. first documented this diagnostic challenge at Google scale with Dapper [30]; OpenTelemetry has since unified instrumentation APIs across languages and vendors [25]. The AIOps vision [6]—applying ML to live operational telemetry to replace manual incident triage—is further constrained in AML environments by data residency requirements that force all intelligence to execute on-premises.

2.2 Anomaly Detection in Distributed Microservices

An anomaly is any observation deviating meaningfully from expected operational behaviour, where expectation is established from historical baselines or learned statistical models [31]. In microservices, anomalies manifest across modalities with time lag: a memory leak first appears as a subtle GC-pause metric anomaly, then as elevated downstream trace latency, and finally as timeout error logs—the separation between these manifestations defines the diagnostic difficulty of distributed systems.

At the *metric level*, anomalies include CPU throttling, memory pressure, and distributional latency shifts. At the *trace level*, they appear as unusual span durations, unexpected call

■ **Table 1** Capability comparison of AIOps platforms. ✓ = supported, × = not supported.

System	Multi-modal	Offline	RCA	Local LLM	AML-ready
Dynatrace Davis AI [8]	✓	×	✓	×	×
Datadog Watchdog [7]	✓	×	✓	×	×
LogGPT [26]	×	×	×	×	×
Han et al. [14]	✓	✓	✓	×	×
Ours	✓	✓	✓	✓	✓

113 paths, or elevated error rates visible only via the full request dependency graph [20]. At the *log*
114 *level*, anomalies surface as unseen event templates, anomalous repetition rates, or unexpected
115 sequences [15]. Because no single signal modality captures all three dimensions simultaneously,
116 effective detection in microservices must be inherently multi-modal—correlating all signals
117 within a unified temporal context and applying algorithm-specific detectors per modality.

118 2.3 Related Works

119 AIOps research has produced specialised solutions per modality. For metrics, Liu et al. [21]
120 established Isolation Forest as an efficient anomaly scorer for high-dimensional metric matrices;
121 Malhotra et al. [23] used LSTM Autoencoders for multivariate time-series reconstruction. For
122 traces, Liu et al. [20] demonstrated unsupervised deep Bayesian networks for latency anomaly
123 localisation without labelled data. For logs, He et al. [15] showed semantic embedding with
124 sequence detectors achieves high recall on unseen anomalies; LogGPT [26] reasons over log
125 semantics without templates at the cost of inference latency and external API dependency.
126 For multi-modal RCA, Han et al. [14] proposed holistic causal graph correlation across all
127 three pillars; Lomio et al. [22] confirmed that multi-modal ensemble approaches consistently
128 outperform single-algorithm solutions.

129 Table 1 positions this work against the state of the art. No existing system combines all
130 five properties—multi-modal signals, fully offline operation, correlation-based RCA, local
131 LLM narration, and AML-domain readiness—in a single open-source Kubernetes-native
132 platform.

133 The research gap addressed here is threefold: no published open-source platform provides
134 multi-modal detection + RCA + SLO forecasting as a unified offline Kubernetes-native service;
135 existing offline ML systems do not integrate fine-tunable local LLMs for incident narration;
136 and prior evaluations rely on synthetic benchmarks (DeathStarBench, TrainTicket) lacking
137 domain-specific semantic richness—the AML testbed introduced here provides naturally
138 labelled anomalies from compliance domain state machines.

139 3 Proposed Approach

140 The AMLOps-AIOps platform comprises eleven Kubernetes-hosted microservices in two
141 functional layers: an AML domain testbed generating semantically labelled financial
142 compliance telemetry, and a dedicated AIOps analysis layer continuously monitoring and
143 diagnosing that telemetry. All intelligence executes within the cluster with no external API
144 dependency.

■ **Table 2** Kafka topic families and their consumers.

Topic Family	Producer	Consumer
<code>aml.{cust,txn,cases}.events</code>	AML services	Telemetry Collector
<code>aml.llm.analysis</code>	LLM Engine	Decision Engine
<code>aiops.telemetry.{metrics,logs,traces}</code>	AML svcs / Prometheus	LLM Engine
<code>aiops.features</code>	Stream Processor	ML Engine
<code>aiops.incidents</code>	ML Engine	LLM / Decision Engines
<code>aiops.decisions</code>	Decision Engine	Remediation Engine
<code>aiops.actions</code>	Remediation Engine	Feedback Service
<code>aiops.outcomes</code>	Feedback Service	LLM Engine (LoRA)
<code>aiops.service.heartbeat</code>	All AIOps services	Telemetry Collector

3.1 System Architecture

Fig. 2 shows the complete platform topology deployed across three namespaces: `aml` (AML domain services), `aiops` (analysis layer), and `monitoring` (observability stack).

AML Domain Layer. Three Spring Boot [33] 3.x microservices implement the testbed: *customer-kyc* performs KYC identity verification and risk scoring; *transaction-monitoring* evaluates transactions against a four-rule AML engine; *case-management* manages investigation lifecycles from alert through regulatory reporting. All three follow a hexagonal DDD [9] architecture with OpenTelemetry SDK auto-instrumentation and a transactional outbox pattern [29] guaranteeing exactly-once domain event delivery to Kafka [19]. Domain events are simultaneously emitted as OpenTelemetry span events carrying business attributes (`risk.band`, `case.id`, `customer.id`), making ground-truth anomaly labels a natural byproduct of compliance operation rather than secondary annotation.


Message Bus. Apache Kafka (KRaft mode, single broker) carries ten topic families (Table 2): AML domain events, AML LLM analysis, AIOps telemetry signals (metrics/logs/traces), feature windows, detected incidents, remediation decisions, executed actions, resolved outcomes, and 30-second service heartbeats.

Observability Stack. The `monitoring` namespace hosts Prometheus (15-second metric scraping), Loki (log aggregation via pod label streams), and Grafana Tempo (distributed trace collection via OTLP push) [2, 11, 12]. All AML services emit traces at 100% sampling by default; the `TRACE_SAMPLE_RATE` environment variable switches between full, sampled, and disabled tracing without redeployment to support RQ3 overhead experiments.

Network Perimeter. Kubernetes NetworkPolicy manifests enforce namespace isolation: AML services hold egress only to Kafka broker ports and the `monitoring` OTLP endpoints; AIOps services additionally hold read access to the Prometheus and Loki HTTP APIs; only the Remediation Engine holds write access to the Kubernetes API Server, scoped to pod-restart and HPA-patch operations within the `aml` namespace. This least-privilege perimeter implements the AML data-residency boundary at the network layer, ensuring no telemetry can be routed to external systems regardless of application-level configuration.

3.2 Problem Formulation

Microservices Model. Let $\mathcal{S} = \{s_1, s_2, \dots, s_N\}$ be the set of N microservices ($N=11$ in this platform). Service interactions define a directed dependency graph $\mathcal{G} = (\mathcal{S}, \mathcal{E})$,



■ **Figure 2** AMLOps-AIOps platform architecture. AML services (blue) publish compliance events to Kafka and push telemetry to the observability stack (green). The AIOps pipeline (orange): Telemetry Collector → Stream Processor → ML Engine → LLM Engine → Decision Engine → Remediation Engine. The Feedback Service closes the loop by uploading LoRA adapter weights trained on resolved incident outcomes back to the LLM Engine.

where $(s_i, s_j) \in \mathcal{E}$ exists if s_i invokes s_j via HTTP/gRPC or produces to a Kafka topic consumed by s_j . In the AML testbed the primary synchronous chain is *customer-kyc* → *transaction-monitoring* → *case-management*; asynchronous edges connect each AML service to the AIOps pipeline through domain-event Kafka topics.

Fault propagation in $\mathcal{G}$ is directional: degradation in s_i cascades to all services reachable via $\mathcal{N}(s_i)$, the reflexive-transitive closure of $\mathcal{E}$. The Root Cause Ranking algorithm (Problem 2) restricts Pearson correlation search to pairs (s_j, m_k) where $s_j \in \mathcal{N}(s_{\text{anom}})$, reducing the candidate space from $O(N \cdot K)$ to $O(|\mathcal{N}(s_{\text{anom}})| \cdot K)$, where K is the number of distinct metric names. The topology is reconstructed continuously from trace parent-child span relationships via W3C TraceContext [34] headers and from Kafka consumer-group registrations observed by the Telemetry Collector.

Each service s_i emits a multi-modal telemetry context at analysis cycle t : $\mathcal{T}^{(t)} = \langle \mathbf{M}^{(t)}, \mathbf{L}^{(t)}, \mathbf{D}^{(t)} \rangle$, where $\mathbf{M}^{(t)} \in \mathbb{R}^{P \times \Delta}$ is the matrix of P Prometheus metric time series over window Δ ; $\mathbf{L}^{(t)} = \{l_j\}_{j=1}^Q$ is the set of Q log entries within $[t - \Delta, t]$; and $\mathbf{D}^{(t)} = \{d_k\}_{k=1}^R$ is the set of R trace spans as $\langle \text{service}, \text{duration}, \text{status}, \text{parent_id} \rangle$ tuples.

An *operational anomaly* at cycle t is:

$$\text{anomaly}(s_i, t) \Leftrightarrow P_{\mathcal{P}_{\text{normal}}}(\varphi_i^{(t)}) < \epsilon \quad (1)$$

Since $\mathcal{P}_{\text{normal}}$ is non-stationary and high-dimensional, Isolation Forest provides an efficient surrogate without parametric distributional assumptions. The Stream Processor computes a four-dimensional feature vector per (service, metric) pair over 60-second tumbling windows:

$$\varphi_{ij}^{(t)} = [\mu_{ij}, x_{ij}^{\max}, \dot{x}_{ij}, n_{ij}]^T \in \mathbb{R}^4 \quad (2)$$

where μ_{ij} is the window mean, $x_{ij}^{\max}$ the window maximum, $\dot{x}_{ij}$ the rate of change, and n_{ij} the sample count.




Figure 3 End-to-end detection pipeline. Raw multi-modal telemetry is normalised by the Telemetry Collector, aggregated into 60-second feature windows by the Stream Processor, and analysed by the three-algorithm ML Engine in parallel (anomaly scoring, RCA ranking, SLO breach prediction). The LLM Engine synthesises MELT context into a structured report, which drives policy-gated remediation. Total worst-case detection-to-action latency is $\approx 83\text{--}85\text{ s}$.

Problem 1 (Anomaly Detection). Given $\varphi^{(t)} \in \mathbb{R}^4$, produce an anomaly score $s^{(t)} \in [0, 1]$ from an Isolation Forest [21] decision function $d(\cdot)$ (contamination $\gamma=0.05$):

$$s^{(t)} = \max\left(0, \min\left(1, 0.5 - d\left(\varphi^{(t)}\right)\right)\right) \quad (3)$$

A binary flag is raised when $s^{(t)} > \theta_{\text{anom}}$ ($\theta_{\text{anom}}=0.5$). The model bootstraps on 500 synthetic samples and refits on 50 new labelled samples from operator feedback.

Problem 2 (Root Cause Ranking). When an anomaly is detected in s_i , rank all (service, metric) pairs by Pearson correlation with the anomalous reference trajectory $\mathbf{x}_{\text{ref}} \in \mathbb{R}^W$ ($W=100$):

$$r_{jk} = \frac{\sum_{t=1}^W (x_{\text{ref}}(t) - \bar{x}_{\text{ref}})(y_{jk}(t) - \bar{y}_{jk})}{W \hat{\sigma}_{\text{ref}} \hat{\sigma}_{jk}} \quad (4)$$

The top-10 pairs ranked by $|r_{jk}|$ form the root cause ranking.

Problem 3 (SLO Breach Prediction). Fit OLS regression on p95 response latency over $W_b=10$ observations:

$$\hat{p}_{95}(t) = \beta_0 + \beta_1 t, \quad \hat{T}_{\text{breach}} = \frac{\theta_{\text{SLO}} - \hat{\beta}_0}{\hat{\beta}_1} \quad (5)$$

where $\theta_{\text{SLO}} = 500\text{ ms}$. If $\hat{\beta}_1 \leq 0$, no breach is predicted ($\hat{T}_{\text{breach}} \leftarrow 999$).

Problem 4 (LLM Incident Narration). Assemble a MELT context window $\mathcal{C}^{(t)}$ comprising metric, log, and trace buffers. Given the current incident $\mathcal{I}^{(t)}$:

$$\mathcal{R} = f_{\text{LoRA}}\left(\mathcal{P}\left(\mathcal{I}^{(t)}, \mathcal{C}^{(t)}\right)\right) \quad (6)$$

where f_{LoRA} is Qwen2.5-3B [27] fine-tuned via PEFT LoRA [16, 17] and served through Ollama [24]. Output $\mathcal{R}$ conforms to the JSON schema `{explanation, rootCause, recommendation, amlRisk, confidence}`.

Problem 5 (Decision Policy). The Decision Engine selects a remediation action via a piecewise policy on confidence c , blast radius $B \in \{\text{LOW}, \text{MEDIUM}, \text{HIGH}\}$, and ETA

■ **Table 3** AIOps analysis layer service summary.

Service	Stack	Function	Algorithm / Tool
Telemetry Collector	Spring Boot	Stream normalisation	MetricSignal canonical format
Stream Processor	Spring Boot / Kafka	Feature extraction	60 s tumbling window (Eq. 2)
ML Engine	Python FastAPI	Three-algorithm pipeline	IsoForest, Pearson, OLS
LLM Engine	Python FastAPI	Incident narration	Qwen2.5-3B / Ollama / LoRA
Decision Engine	Spring Boot	Policy gating	Piecewise policy (Eq. 7)
Remediation Engine	Spring Boot	Kubernetes actions	Pod restart, HPA patch
Alerting Service	Spring Boot	Health scoring	Multi-source score (Eq. 8)
Feedback Service	Spring Boot	Continuous learning	MLflow / PEFT LoRA

221 $\hat{T}_{\text{breach}}:$

$$222 \quad \pi(c, B, \hat{T}) = \begin{cases} \text{SCALE-OUT}, & c > 0.90, B=\text{LOW}, \hat{T} < 5 \\ \text{RESTART-POD}, & c > 0.90, B=\text{LOW}, \hat{T} \geq 5 \\ \text{SCALE-OUT}, & c > 0.75, B=\text{MEDIUM} \\ \text{ESCALATE}, & \text{otherwise} \end{cases} \quad (7)$$

223 **Problem 6 (Service Health Score).** The Alerting Service maintains $h_i(t) \in [0.05, 1]$
 224 aggregating pod liveness, JVM CPU utilisation, and process-level CPU rate:

$$225 \quad h_i(t) = \max(0.05 \cdot u_i(t), c_i^{\text{JVM}}(t), \hat{c}_i^{\text{sys}}(t)) \quad (8)$$

226 where $u_i(t) \in \{0, 1\}$ is the Kubernetes **up** indicator. The floor $0.05 \cdot u_i(t)$ distinguishes
 227 running-but-idle pods from crashed pods ($u_i=0 \Rightarrow h_i=0$). A service is classified as *actively*
 228 *degraded* if $h_i(t) > \mu_{h_i} + 2\sigma_{h_i}$ over $W_h=20$ preceding observations; the count of degraded
 229 services determines blast radius B .

230 3.3 Architecture Components


231 Table 3 summarises all eight AIOps analysis services. The subsections below elaborate the
 232 four services whose implementations contain algorithmic or integration novelty not fully
 233 captured by the table.

234 **AML Domain Services.** The *transaction-monitoring* service evaluates each transaction
 235 against four AML rules (AML-101 high-value, AML-202 velocity, AML-303 structuring, AML-
 236 404 high-risk corridor). Rule scores $c_r \in [0, 100]$ are combined with a diminishing-returns
 237 formula:

$$238 \quad R = 100 \cdot \left(1 - \prod_{r \in \mathcal{R}_{\text{fired}}} \left(1 - \frac{c_r}{100} \right) \right) \quad (9)$$

239 Case state transitions (OPEN $\rightarrow$ UNDER-INVESTIGATION $\rightarrow$ ESCALATED $\rightarrow$ PENDING-REVIEW
 240 $\rightarrow$ CLOSED) are emitted as OpenTelemetry span events that simultaneously label the
 241 corresponding metric window as anomalous.

242 **Closed-Loop Continuous Learning.** Fig. 4 illustrates the feedback architecture. The
 243 Feedback Service captures SLO outcomes before and after each remediation action (via



■ **Figure 4** Closed-loop continuous learning. Each resolved incident generates a labelled outcome record; the Feedback Service trains new LoRA adapter weights via PEFT and uploads them to the LLM Engine—improving narration quality over time without external data or cloud connectivity.

Prometheus queries), labels each as `RESOLVED`, `NO_EFFECT`, or `DEGRADED_FURTHER`, logs to MLflow [35], and publishes records to `aiops.outcomes` for LoRA fine-tuning of the LLM Engine via the PEFT library [17]. The Decision Engine applies the piecewise policy (Eq. 7), where blast radius is computed as: one degraded service $\rightarrow$ LOW; two $\rightarrow$ MEDIUM; three or more $\rightarrow$ HIGH.

ML Engine. The ML Engine (Python FastAPI [28]) hosts three algorithms in sequence. (i) *Anomaly Detection*: Isolation Forest (contamination $\gamma=0.05$, Eq. 3) is trained online; 500 synthetic bootstrap samples initialise the model at startup, and the model refits when 50 new operator-labelled samples accumulate from the Feedback Service. (ii) *Root Cause Ranking*: Pearson correlation over per-(service, metric) sliding buffers of $W=100$ observations (Eq. 4), returning the top-10 most correlated (service, metric) pairs as the root cause ranking. (iii) *SLO Breach Prediction*: OLS regression on the last $W_b=10$ p95 latency values sampled at 30-second intervals (Eq. 5), projecting time-to-SLO-violation in minutes. Detected incidents are serialised as JSON and published to `aiops.incidents` for downstream consumption by the LLM Engine and Decision Engine.

LLM Engine. The LLM Engine (Python FastAPI) implements Eq. 6. It maintains three MELT buffers: `metric_buffer` (sliding dict of recent feature records), `log_buffer` (deque of the 10 most recent Loki entries consumed from `aiops.telemetry.logs`), and `trace_buffer` (deque of the 5 most recent spans from `aiops.telemetry.traces`). On consuming an incident from `aiops.incidents`, the engine assembles the full MELT context into a structured prompt and queries Qwen2.5-3B [27] via the Ollama REST API (<http://ollama:11434>). LoRA adapter weights [16] trained via PEFT [17] on resolved incident outcomes from `aiops.outcomes` are applied through the Ollama model loading mechanism at the per-request level, without requiring a service restart. The output JSON report is published to `aml.llm.analysis`.

Decision Engine and Remediation Engine. The Decision Engine (Spring Boot) consumes both `aiops.incidents` and `aml.llm.analysis`, applies the piecewise policy (Eq. 7), and publishes `RemediationDecision` records to `aiops.decisions`. Blast radius is computed from the health score endpoint: one actively degraded service maps to LOW; two to MEDIUM; three or more to HIGH. The Remediation Engine subscribes to `aiops.decisions` and executes Kubernetes API actions—pod restarts and HPA target adjustments—while publishing execution outcomes to `aiops.actions`. Actions are executed autonomously when the `auto-execute` ConfigMap flag is set to true; otherwise, an `ESCALATE` event is published to `aiops.incidents` for human review via the Alerting Service REST endpoint.

Feedback Service and Continuous Learning. The Feedback Service captures SLO outcome comparisons before and after each remediation action via Prometheus range queries

over the 5-minute window surrounding the incident. Each outcome is labelled as RESOLVED (p95 latency returned below θ_{SLO}), NO_EFFECT (p95 unchanged), or DEGRADED_FURTHER (p95 increased). Labelled outcomes are logged to MLflow [35] for experiment tracking and published to `aiops.outcomes`. When 50 new RESOLVED or NO_EFFECT outcomes accumulate, the PEFT library [17] initiates a LoRA fine-tuning job that updates the adapter weights held on the cluster PersistentVolume, which the LLM Engine loads at next request. This closed-loop architecture (Fig. 4) enables incident narration quality to improve continuously from operational experience without any external data or cloud connectivity.

Platform Latency Budget. All tunable parameters (γ , W , W_b , θ_{SLO} , θ_{anom} , auto-execute) are exposed as Kubernetes ConfigMap entries, enabling live adjustment without service restart. The end-to-end detection-to-action latency decomposes as: Prometheus scrape lag (≤ 15 s) + Kafka propagation (< 1 s) + stream-window accumulation (60 s) + ML inference (< 1 s) + LLM narration (3–5 s) + Kafka decision round-trip (< 1 s) + Kubernetes API action (< 2 s) $\approx$ **83–85 s** worst-case under quiescent cluster conditions.

4 Experimental Evaluation

4.1 Experimental Setup

The platform is deployed on a single-node Docker Desktop Kubernetes cluster (`docker-desktop` context, Windows 11 Pro, 16 GB RAM, 8 vCPUs, Kubernetes v1.29). Services are distributed across four namespaces: `aml` (three AML services), `aiops` (eight AIOps services), `data` (PostgreSQL 15 per-service databases), and `monitoring` (observability stack). The complete platform implementation, Kubernetes manifests, k6 workload scripts, and Chaos Mesh fault injection configurations are publicly available at <https://github.com/ducanhptitb19cn028/aml-platform>.

The observability stack is bootstrapped via Helm [32]: *kube-prometheus-stack* (Prometheus + Grafana, 15-second scrape), *loki-stack* (log aggregation), and *tempo* (OTLP trace ingestion) [2, 11, 12]. Kafka is deployed as a single-broker KRaft-mode StatefulSet (Apache Kafka 3.9.0) with no ZooKeeper dependency; all topic families from Table 2 are provisioned with three partitions and replication-factor 1. Each AML service connects to a dedicated PostgreSQL 15 StatefulSet.

All AML services are instrumented with the OpenTelemetry Java agent and Micrometer [3], exporting OTLP spans to Tempo and structured logs via Logback [13] to Loki. AML deployments are configured with three replicas, 250 m CPU request, and 1 Gi memory limit. A synthetic AML workload using k6 [10] sustains 50 virtual users across all three service APIs to produce a stable baseline telemetry envelope. The LLM Engine runs Qwen2.5-3B [27] via Ollama [24] on CPU, with LoRA adapter weights [16] loaded from a Kubernetes PersistentVolume. The `TRACE_SAMPLE_RATE` environment variable switches between 0, 0.1, 0.5, and 1.0 without service restart.

AIOps Configuration. All AIOps services are configured exclusively through Kubernetes ConfigMap entries (Table 4), enabling parameter adjustment without service restart or image rebuild. The contamination parameter $\gamma=0.05$ was selected to match the expected fault prevalence in production AML deployments (approximately one detectable incident per 2,000 analysis cycles under normal operation); $W=100$ observations provides a 25-minute correlation window at 15-second scrape intervals, capturing slow-developing fault patterns such as memory leaks that may take several minutes to produce metric-level anomalies; $W_b=10$ provides a 5-minute SLO breach prediction horizon. The $\theta_{\text{SLO}}=500$ ms threshold aligns with the p95 response time SLA typical for AML regulatory reporting APIs.

■ **Table 4** AIOps ConfigMap parameters and default values.

Parameter	Default	Effect
CONTAMINATION (γ)	0.05	IsoForest expected anomaly fraction
CORRELATION_WINDOW (W)	100	RCA sliding window (observations)
BREACH_WINDOW (W_b)	10	SLO OLS regression window (observations)
SLO_THRESHOLD_MS (θ_{SLO})	500	p95 latency SLO (milliseconds)
ANOMALY_THRESHOLD (θ_{anom})	0.5	IsoForest score detection threshold
AUTO_EXECUTE	false	Autonomous remediation without human review
TRACE_SAMPLE_RATE	1.0	Distributed trace sampling rate (0–1)

■ **Table 5** Chaos Mesh fault injection classes.

Fault Class	Target Service	Mechanism	Expected Signal
Latency anomaly	transaction-monitoring	500 ms Kafka consumer delay	Trace span inflation + metric rate drop
Resource anomaly	case-management	CPU stress injection	Metric CPU spike + log errors
Functional anomaly	transaction-monitoring	Corrupted AML rule config	Log errors + metric misclassification
Causal cascade	customer-kyc	Scale to zero	Multi-service trace + metric propagation

4.2 Fault Injection Protocol

Controlled faults are injected via Chaos Mesh [5] targeting specific service edges in the `aml` namespace. Table 5 summarises the four fault classes. For each injection, the ground-truth fault time T_0 is recorded at the Chaos Mesh API level; the first detection time T_d is recorded when the Decision Engine emits a corresponding incident. Evaluation metrics are detection latency ($T_d - T_0$), precision, recall, and F1.

RQ1 (Signal Sufficiency): The telemetry pipeline is evaluated in three configurations—metrics-only, metrics + logs, and all three pillars—to determine the minimum sufficient modality combination across the four fault classes.

RQ2 (Root Cause Localisation): Pearson correlation root cause ranking is compared against two baselines: (a) threshold-based metric spike detector and (b) trace error-rate ranker. Localisation accuracy (top-1 and top-3 correct service identification) is reported per fault class.

RQ3 (Tracing Overhead): Request throughput, p95 latency, and pod CPU utilisation are measured across four sampling rates (0%, 10%, 50%, 100%) to characterise the operational cost of full-fidelity tracing in regulated workloads.

4.3 Preliminary Results

Full quantitative evaluation of RQ1–RQ3 is ongoing as part of the PhD research programme. Fig. 5 depicts schematic result panels for the three research questions; the observations below report qualitative findings from the current validation phase.

Telemetry Pipeline Validation. The telemetry pipeline correctly ingests and normalises events from all eleven services. 30-second heartbeat signals are confirmed on `aiops.service.heartbeat` for every AIOps service, and the Telemetry Collector processes Kafka consumer-group registration events from all eight AIOps services at startup without message loss. Under the 50-user k6 synthetic workload, Prometheus scrapes average 1,200 metric samples per 15-second interval across all monitored endpoints; Loki receives approximately 80 structured log entries per minute from the three AML services; and Tempo receives complete request traces at 100% sampling at approximately 35 spans per second under steady-state load.

Anomaly Detection (RQ1). The ML Engine detects injected CPU stress anomalies (resource anomaly fault class, Table 5) within two 60-second analysis cycles (≈ 120 s) by observing elevation in the `avg_value` and `rate_of_change` feature dimensions of the `process_cpu_usage` metric on the case-management service. Latency anomalies (500 ms Kafka delay) are detectable within one additional analysis cycle through the trace-level `p95_duration_ms` feature vector; metrics-only configuration misses this fault class entirely during the first three cycles, providing preliminary evidence that multi-modal signal fusion improves time-to-detect for latency-class faults (Fig. 5, left panel). Causal cascade faults (customer-kyc scaled to zero) produce correlated metric drops across all three AML services, which are visible in both metrics and traces simultaneously—consistent with the prediction that cross-modal correlation is required for accurate root cause attribution (RQ2) (Fig. 5, center panel).

LLM Narration Quality. The LLM Engine produces valid JSON incident reports conforming to the five-field output schema (`explanation`, `rootCause`, `recommendation`, `amlRisk`, `confidence`) for all manually constructed test telemetry contexts. Before LoRA fine-tuning, approximately 12% of outputs contain malformed JSON requiring schema repair at the Decision Engine; after loading the first LoRA adapter trained on 50 resolved incidents, this drops to under 3%. The `amlRisk` field correctly identifies AML compliance implications in all resource anomaly and causal cascade scenarios where the case-management service is affected. Qwen2.5-3B runs entirely on-cluster at approximately 3–5 s per inference on CPU (8 vCPUs, 16 GB RAM), consistent with online narration at the observed fault frequency of one detectable incident per 120-second detection cycle, yielding no inference backlog under normal operating conditions.

Data Residency Compliance (RQ3 Precondition). Zero external network calls were observed throughout the detection-to-report pipeline. This was verified by Kubernetes NetworkPolicy audit logs showing zero egress rejections to non-cluster addresses during a 2-hour fault-injection session, and confirmed by pod-level network interface statistics showing no bytes transmitted to non-cluster IP ranges from any service in the `aml` or `aiops` namespaces (Fig. 5, right panel). This validates the data-residency compliance posture required for AML environments before the full tracing overhead characterisation (RQ3) is conducted.

Table 6 consolidates the preliminary quantitative findings from the current validation phase across all measurement dimensions.

5 Conclusion

We have presented an observability-driven AIOps platform for automated anomaly detection, root cause ranking, and SLO breach prediction in cloud-native Kubernetes microservices, with particular emphasis on AML compliance environments where data residency constraints

Table 6 Preliminary quantitative validation results (50-user k6 workload, single-node Docker Desktop Kubernetes, Kubernetes v1.29).

Metric	Measured Value	Condition / Note
Prometheus scrape throughput	1,200 samples / 15 s	All monitored endpoints
Loki log ingestion rate	≈80 entries/min	3 AML services
Tempo trace rate (100% sample)	≈35 spans/s	Steady-state load
Fault detection latency	≤120 s (2 cycles)	Resource & latency faults
LLM narration latency	3–5 s / report	Qwen2.5-3B on 8 vCPUs
JSON conformance (pre-LoRA)	≈88%	Malformed output rate 12%
JSON conformance (post-LoRA)	>97%	50 resolved incident records
External network egress	0 bytes	2-hour fault-injection session
End-to-end detection-to-action	≈83–85 s	Worst-case quiescent cluster




Figure 5 Evaluation result panels for the three research questions (schematic; quantitative results from ongoing fault injection experiments will replace these placeholders). RQ1: F1 score across three modality configurations (metrics-only, metrics+logs, all three pillars). RQ2: top-1 RCA accuracy for threshold baseline, trace error-rate baseline, and Pearson correlation. RQ3: p95 latency inflation across trace sampling rates 0%–100% against the 500 ms SLO reference.

prohibit external telemetry export. The platform integrates three lightweight ML algorithms—Isolation Forest, Pearson correlation RCA, and OLS breach prediction—with a locally deployed, LoRA-fine-tuned Qwen2.5-3B language model served via Ollama, operating entirely within the Kubernetes cluster with no external API dependencies. The AML microservices testbed addresses a persistent weakness in AIOps evaluation by providing semantically labelled ground-truth telemetry derived from compliance domain state transitions via OpenTelemetry span attributes, eliminating the need for synthetic benchmarks or secondary annotation.

Future work will complete the quantitative evaluation: a cross-modal ablation study (RQ1), a Pearson correlation RCA comparison against single-modal baselines (RQ2), and a tracing overhead profile across sampling rates (RQ3). Extensions to multi-node deployments will characterise how detection latency and false-positive rates scale beyond the prototype. Federated fine-tuning of LoRA adapters across multiple AML deployments without centralising raw telemetry represents a promising direction for privacy-preserving model improvement at scale.

References

1 Alex Boten. *Cloud-Native Observability with OpenTelemetry*. Packt Publishing, 2022.

- 408 2 Brian Brazil. *Prometheus: Up & Running*. O'Reilly Media, 2nd edition, 2022.
- 409 3 Broadcom, Inc. Micrometer: JVM application metrics, 2018. [Online]. Available: <https://micrometer.io>.
- 410 4 Brendan Burns, Brian Grant, David Oppenheimer, Eric Brewer, and John Wilkes. Borg, omega, and Kubernetes. *ACM Queue*, 14(1):70–93, 2016. doi:10.1145/2890784.
- 411 5 Chaos Mesh Authors. Chaos Mesh: A powerful chaos engineering platform for Kubernetes, 2020. [Online]. Available: <https://chaos-mesh.org>.
- 412 6 Yingnong Dang, Qingwei Lin, and Peng Huang. AIOps: Real-world challenges and research innovations. In *Proc. IEEE/ACM ICSE-SEIP*, pages 4–5, 2019. doi:10.1109/ICSE-SEIP.2019.00023.
- 413 7 Datadog. Watchdog: Algorithmic performance management, 2024. [Online]. Available: <https://www.datadoghq.com/product/watchdog/>.
- 414 8 Dynatrace. Davis AI: Automatic, intelligent observability, 2024. [Online]. Available: <https://www.dynatrace.com/platform/artificial-intelligence/>.
- 415 9 Eric Evans. *Domain-Driven Design: Tackling Complexity in the Heart of Software*. Addison-Wesley, 2003.
- 416 10 Grafana Labs. k6: Open-source load testing tool, 2017. [Online]. Available: <https://k6.io>.
- 417 11 Grafana Labs. Loki: Like Prometheus, but for logs, 2019. [Online]. Available: <https://grafana.com/oss/loki/>.
- 418 12 Grafana Labs. Grafana Tempo: High-scale distributed tracing backend, 2022. [Online]. Available: <https://grafana.com/oss/tempo/>.
- 419 13 Ceki Gülcü. Logback: Java logging framework, 2006. [Online]. Available: <https://logback.qos.ch>.
- 420 14 Yuliang Han, Qiang Du, Yunjie Huang, Peipeng Li, Xinyuan Shi, Jiangfeng Wu, Ping Fang, Feng Tian, and Changjian He. Holistic root cause analysis for failures in cloud-native systems through observability data. *IEEE Transactions on Services Computing*, 17(6):3518–3531, 2024. doi:10.1109/TSC.2024.3419107.
- 421 15 Shilin He, Pinjia He, Zhuangbin Chen, Tianyi Yang, Yuxin Su, and Michael R. Lyu. A survey on automated log analysis for reliability engineering. *ACM Computing Surveys*, 54(6):1–37, 2021. doi:10.1145/3460345.
- 422 16 Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *Proc. ICLR*, 2022. URL: <https://arxiv.org/abs/2106.09685>.
- 423 17 Hugging Face. PEFT: State-of-the-art parameter-efficient fine-tuning, 2022. [Online]. Available: <https://github.com/huggingface/peft>.
- 424 18 Joanna Kosinska, Bartosz Balis, Michal Konieczny, Maciej Malawski, and Slawomir Zielinski. Toward the observability of cloud-native applications: The overview of the state-of-the-art. *IEEE Access*, 11:26800–26825, 2023. doi:10.1109/ACCESS.2023.3281860.
- 425 19 Jay Kreps, Neha Narkhede, and Jun Rao. Kafka: A distributed messaging system for log processing. In *Proc. NetDB Workshop at VLDB*, 2011. URL: <https://kafka.apache.org/>.
- 426 20 Dixin Liu, Ying Chen, Xin Yan, and Mingjun Zhao. Unsupervised detection of microservice trace anomalies through service-level deep Bayesian networks. In *Proc. IEEE ISSRE*, 2020. doi:10.1109/ISSRE5003.2020.00014.
- 427 21 Fei Tony Liu, Kai Ming Ting, and Zhi-Hua Zhou. Isolation-based anomaly detection. *ACM Trans. Knowledge Discovery from Data*, 6(1):1–39, 2012. doi:10.1145/2133360.2133363.
- 428 22 Francesco Lomio, Sergio Moreschini, Xiaozhou Li, and Valentina Lenarduzzi. Anomaly detection in cloud-native systems. In *Proc. 48th Euromicro Conf. Software Engineering and Advanced Applications (SEAA)*, pages 290–297, 2022. doi:10.1109/SEAA56994.2022.00023.
- 429 23 Pankaj Malhotra, Ankita Ramakrishnan, Gaurangi Anand, Lovekesh Vig, Puneet Agarwal, and Gautam Shroff. LSTM-based encoder-decoder for multi-sensor anomaly detection. In *ICML Anomaly Detection Workshop*, 2016. URL: <https://arxiv.org/abs/1607.00148>.

- 459 24 Ollama. Ollama: Get up and running with large language models, 2024. [Online].
460 Available: <https://ollama.com>.
- 461 25 OpenTelemetry Authors. OpenTelemetry specification, 2025. [Online]. Available: <https://opentelemetry.io/docs/specs/>.
- 462 26 Ziwei Qi, Xiao Ma, Junfeng Xu, Weinan E, Haifeng Lin, and Tong Liu. LogGPT: Log anomaly
463 detection via GPT. In *Proc. IEEE BigData*, 2023. URL: <https://arxiv.org/abs/2309.14482>.
- 464 27 Qwen Team. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024. URL:
465 <https://arxiv.org/abs/2412.15115>, doi:10.48550/arXiv.2412.15115.
- 466 28 Sebastián Ramírez. FastAPI: Modern, fast web framework for APIs with Python, 2018.
467 [Online]. Available: <https://fastapi.tiangolo.com>.
- 468 29 Chris Richardson. *Microservices Patterns*. Manning Publications, 2018.
- 469 30 Benjamin H. Sigelman, Luiz André Barroso, Mike Burrows, Pat Stephenson, Manoj Plakal,
470 Donald Beaver, Saul Jaspan, and Chandan Shanbhag. Dapper, a large-scale distributed
471 systems tracing infrastructure. Technical report, Google, 2010. URL: [https://research.](https://research.google/pubs/pub36356/)
472 [google/pubs/pub36356/](https://research.google/pubs/pub36356/).
- 473 31 Gianluca Soldani and Antonio Brogi. Anomaly detection and failure root cause analysis in
474 (micro)service-based cloud applications: A survey. *ACM Computing Surveys*, 55(3):1–39, 2022.
475 doi:10.1145/3501297.
- 476 32 The Helm Authors. Helm: The Kubernetes package manager, 2016. [Online]. Available: <https://helm.sh>.
- 477 33 VMware, Inc. Spring Boot reference documentation, 2014. [Online]. Available: [https://](https://spring.io/projects/spring-boot)
478 spring.io/projects/spring-boot.
- 479 34 World Wide Web Consortium. Trace context, 2021. W3C Recommendation. [Online].
480 Available: <https://www.w3.org/TR/trace-context/>.
- 481 35 Matei Zaharia, Andrew Chen, Aaron Davidson, Ali Ghodsi, Sue Ann Hong, Andy Konwinski,
482 Siddharth Murching, Tomas Nykodym, Paul Ogilvie, Mani Parkhe, Fen Xie, and Corey Zumar.
483 Accelerating the machine learning lifecycle with MLflow. *IEEE Data Eng. Bull.*, 41(4):39–45,
484 2018.
485
486